



中山大學

SUN YAT-SEN UNIVERSITY

《操作系统》 Assignment3

年冬冬

18364073

2021 年 1 月 23 日

目录

1. xv6 Lab: Multithreading/Uthread: switching between threads	3
1.1 题目要求	3
1.2 实验代码	3
1.3 实验结果	5
2. Xv6 lab: Lock/Memory allocator.....	6
2.1 题目要求	6
2.2 实验代码	6
2.3 实验结果	8
3. Xv6 lab: Lock/Buffer cache	10
3.1 题目要求	10
3.2 实验代码	10
3.3 实验结果	14
4. Xv6 lab: File System/Large files.....	15
4.1 题目要求	15
4.2 实验代码	15
4.3 实验结果	18

1.xv6 Lab: Multithreading/Uthread: switching between threads

1.1 题目要求

在本练习中，你将为用户级线程系统设计上下文切换机制，然后实现它。首先，你的 xv6 有两个文件 `user/uthread.c` 和 `user/uthread_switch.S`，以及 `Makefile` 中的一个用于构建 `uthread` 程序的规则。`uthread.c` 包含大多数用户级线程包，以及三个简单测试线程的代码。线程包缺少一些用于创建线程和在线程之间切换的代码。你的工作是想出一个创建线程的计划，并保存/恢复寄存器以在线程之间切换，并实施该计划。

1.2 实验代码

1. 在 `struct thread` 添加寄存器字段，能够保存进程内容。其中 `ra` 为返回地址，`sp` 为栈顶指针地址。

```
struct thread {
    /*stored registers*/
    uint64 ra;
    uint64 sp;
    // callee-saved
    uint64 s0;
    uint64 s1;
    uint64 s2;
    uint64 s3;
    uint64 s4;
    uint64 s5;
    uint64 s6;
    uint64 s7;
    uint64 s8;
    uint64 s9;
    uint64 s10;
    uint64 s11;

    char    stack[STACK_SIZE]; /* the thread's stack */
    int     state;              /* FREE, RUNNING, RUNNABLE */
};
```


2. 在 `thread_create()` 中添加线程进程入口，设置栈顶指针 `sp` 进行用户内存分配。

```
void
thread_create(void (*func)())
{
    struct thread *t;

    for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
        if (t->state == FREE) break;
    }
    t->state = RUNNABLE;
    // YOUR CODE HERE
    t->ra=(uint64)func;
    t->sp=(uint64)&t->stack[STACK_SIZE-1];
}
```

3. 在 `thread_schedule()` 中添加调用

```

Open ▾  *uthread.c
~/xv6-labs-2020/user

next_thread = 0;
t = current_thread + 1;
for(int i = 0; i < MAX_THREAD; i++){
    if(t >= all_thread + MAX_THREAD)
        t = all_thread;
    if(t->state == RUNNABLE) {
        next_thread = t;
        break;
    }
    t = t + 1;
}

if (next_thread == 0) {
    printf("thread_schedule: no runnable threads\n");
    exit(-1);
}

if (current_thread != next_thread) {          /* switch threads? */
    next_thread->state = RUNNING;
    t = current_thread;
    current_thread = next_thread;
    /* YOUR CODE HERE
     * Invoke thread_switch to switch from t to next_thread:*/
    thread_switch((uint64)t,(uint64)next_thread);
} else
    next_thread = 0;
}

```

4. 在 uthread_switch.S 中添加寄存器

```

thread_switch:
    /* YOUR CODE HERE */
    sd ra, 0(a0)
    sd sp, 8(a0)
    sd s0, 16(a0)
    sd s1, 24(a0)
    sd s2, 32(a0)
    sd s3, 40(a0)
    sd s4, 48(a0)
    sd s5, 56(a0)
    sd s6, 64(a0)
    sd s7, 72(a0)
    sd s8, 80(a0)
    sd s9, 88(a0)
    sd s10, 96(a0)
    sd s11, 104(a0)

    ld ra, 0(a1)
    ld sp, 8(a1)
    ld s0, 16(a1)
    ld s1, 24(a1)
    ld s2, 32(a1)
    ld s3, 40(a1)
    ld s4, 48(a1)
    ld s5, 56(a1)
    ld s6, 64(a1)
    ld s7, 72(a1)
    ld s8, 80(a1)
    ld s9, 88(a1)
    ld s10, 96(a1)
    ld s11, 104(a1)
    ret    /* return to ra */

```

1.3 实验结果

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help

xv6 kernel is booting

hart 1 starting
hart 2 starting
init: starting sh
$ uthread
thread_a started
thread_b started
thread_c started
thread_a 0
thread_b 0
thread_c 1
thread_a 1
thread_b 1
thread_c 2
thread_a 2
thread_b 2
thread_c 3
thread_a 3
thread_b 3
thread_c 4
thread_a 4
```

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help

thread_b 93
thread_c 94
thread_a 94
thread_b 94
thread_c 95
thread_a 95
thread_b 95
thread_c 96
thread_a 96
thread_b 96
thread_c 97
thread_a 97
thread_b 97
thread_c 98
thread_a 98
thread_b 98
thread_c 99
thread_a 99
thread_b 99
thread_c: exit after 100
thread_a: exit after 100
thread_b: exit after 100
thread_schedule: no runnable threads
$
```

2.Xv6 lab: Lock/Memory allocator

2.1 题目要求

对于每个锁，acquire 都会维护获取该锁的调用次数，以及获取中的循环尝试但未设置锁的次数。如果存在锁争用，则获取循环的迭代次数将很大。kalloc_test 调用一个系统调用，系统调用返回 kmem 和 bcache 锁的循环迭代总数。

kalloc_test 中锁争用的根本原因是 kalloc() 具有单个空闲列表，并受单个锁保护。要删除锁争用，需要重新设计内存分配器以避免单个锁和列表。基本思想是为每个 CPU 维护一个空闲列表，每个列表都有自己的锁。不同 CPU 上的分配和释放可以并行运行，因为每个 CPU 将在不同的列表上运行。主要的挑战将是处理一个 CPU 的空闲列表为空，而另一个 CPU 的列表具有空闲内存的情况。在这种情况下，一个 CPU 必须“窃取”另一 CPU 空闲列表的一部分，窃取可能会引入锁争用。

2.2 实验代码

1. 首先将 kmem 修改为数组，这样每个 cpu 对应一个 freelist 和 lock。

```
struct {
    struct spinlock lock;
    struct run *freelist;
} kmem[NCPU];
```

2. 修改 kernel/kalloc.c 中的 kinit 函数，初始化 kmem 时将每个 cpu 对应 kmem[i] 都初始化。

```
void
kinit()
{
    for(int i=0; i<NCPU; i++){
        initlock(&kmem[i].lock, "kmem");
    }
    freerange(end, (void*)PHYSTOP);
}
```

3. 修改 kfree 函数

```
1. void
2. kfree(void *pa)
3. {
4.     struct run *r;
5.
6.     if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
7.         panic("kfree");
8.     // Fill with junk to catch dangling refs.
9.     memset(pa, 1, PGSIZE);
10.
11.     r = (struct run*)pa;
12.
```

```

13.  push_off();// 关中断
14.  int i=cuid();// CPU id
15.  acquire(&kmem[i].lock);
16.  r->next = kmem[i].freelist;
17.  kmem[i].freelist = r;
18.  release(&kmem[i].lock);
19.
20.  pop_off();//开中断
21. }

```

4. 修改 kalloc 函数

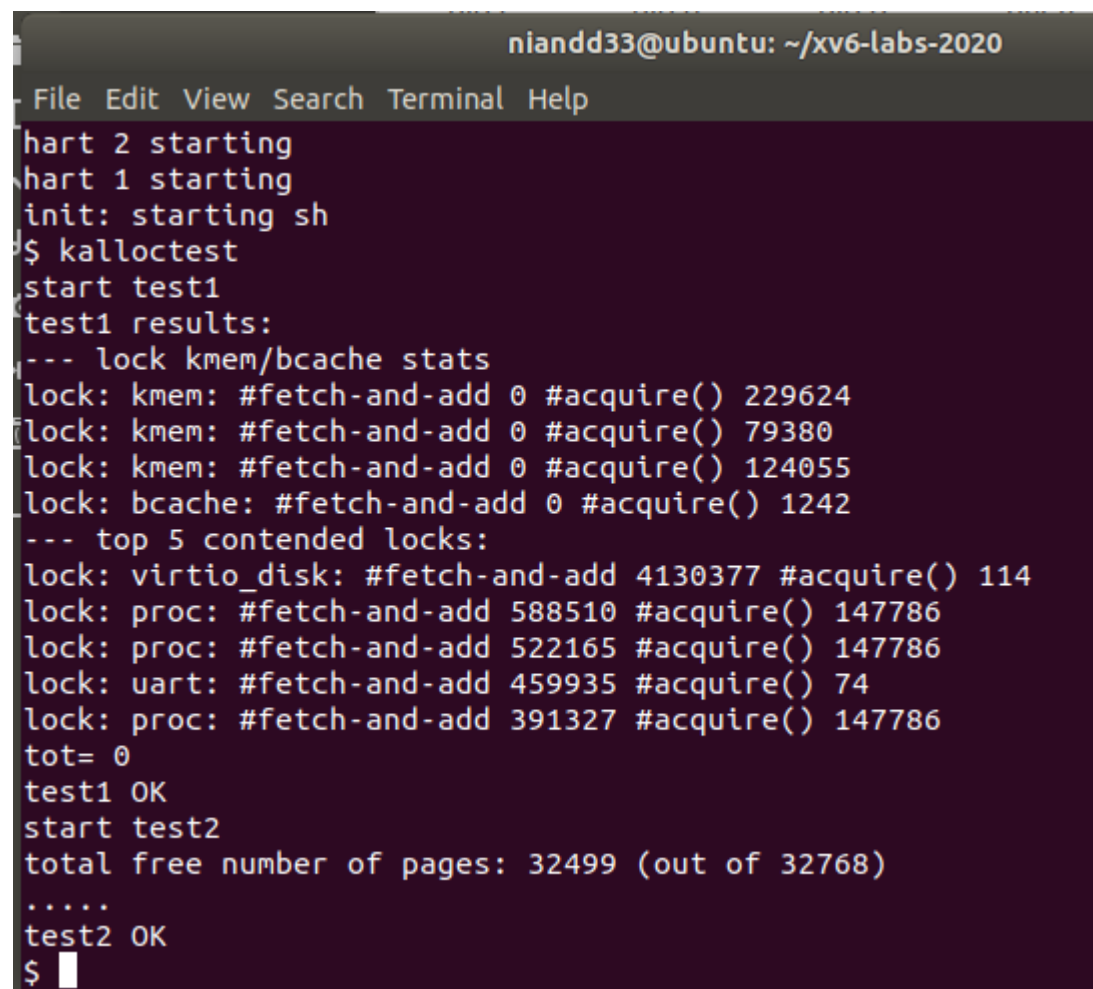
```

1.  void *
2.  kalloc(void)
3.  {
4.      struct run *r;
5.
6.      push_off();// 关中断
7.      int i=cuid();// CPU id
8.      acquire(&kmem[i].lock);
9.      r = kmem[i].freelist;
10.     if(r)
11.         kmem[i].freelist = r->next;
12.     release(&kmem[i].lock);
13.
14.     if(!r)//如果当前 cpu 对应 freelist 为空
15.     {
16.         for(int j=0;j<NCPU;j++)//从其他 cpu 的 freelist 中借用,遍历所有 CPU
17.         {
18.             if(j!=i)
19.             {
20.                 acquire(&kmem[j].lock);//获取该 CPU 的锁
21.                 if(kmem[j].freelist)
22.                 {
23.                     r=kmem[j].freelist;//获取该 CPU 的空闲列表
24.                     kmem[j].freelist=r->next; //该 CPU 空闲列表-1, 给窃取者
25.                     release(&kmem[j].lock);
26.                     break;
27.                 }
28.                 release(&kmem[j].lock);
29.             }
30.         }
31.     }

```

```
32.  
33. pop_off();//开中断  
34.  
35. if(r)  
36.     memset((char*)r, 5, PGSIZE); // fill with junk  
37. return (void*)r;  
38. }
```

2.3 实验结果

A terminal window titled 'niandd33@ubuntu: ~/xv6-labs-2020' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal output shows the progress of xv6 booting and running tests. It starts with 'hart 2 starting', 'hart 1 starting', and 'init: starting sh'. Then, the 'kalloc test' is run, showing 'start test1' and 'test1 results:'. The results include lock statistics for kmem, bcache, and virtio_disk, and a list of top 5 contended locks. After 'test1 OK', 'start test2' is run, showing 'total free number of pages: 32499 (out of 32768)' and 'test2 OK'. The prompt '\$' is visible at the bottom.

```
niandd33@ubuntu: ~/xv6-labs-2020  
File Edit View Search Terminal Help  
hart 2 starting  
hart 1 starting  
init: starting sh  
$ kalloc test  
start test1  
test1 results:  
--- lock kmem/bcache stats  
lock: kmem: #fetch-and-add 0 #acquire() 229624  
lock: kmem: #fetch-and-add 0 #acquire() 79380  
lock: kmem: #fetch-and-add 0 #acquire() 124055  
lock: bcache: #fetch-and-add 0 #acquire() 1242  
--- top 5 contended locks:  
lock: virtio_disk: #fetch-and-add 4130377 #acquire() 114  
lock: proc: #fetch-and-add 588510 #acquire() 147786  
lock: proc: #fetch-and-add 522165 #acquire() 147786  
lock: uart: #fetch-and-add 459935 #acquire() 74  
lock: proc: #fetch-and-add 391327 #acquire() 147786  
tot= 0  
test1 OK  
start test2  
total free number of pages: 32499 (out of 32768)  
.....  
test2 OK  
$
```



```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
.....
test2 OK
$ usertests sbrkmuch
usertests starting
test sbrkmuch: OK
ALL TESTS PASSED
$ usertests
usertests starting
test manywrites: OK
test execout: OK
test copyin: OK
test copyout: OK
test copyinstr1: OK
test copyinstr2: OK
test copyinstr3: OK
test rwsbrk: OK
test truncate1: OK
test truncate2: OK
test truncate3: OK
test reparent2: OK
test pgbug: OK
test sbrkbugs: usertrap(): unexpected scause 0x000000000000000c pid=3247
                sepc=0x000000000000056a0 stval=0x000000000000056a0
usertrap(): unexpected scause 0x000000000000000c pid=3248
```

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
test validatetest: OK
test stacktest: usertrap(): unexpected scause 0x000000000000000d pid=6283
                sepc=0x000000000000022cc stval=0x000000000000fb90
OK
test opentest: OK
test writetest: OK
test writebig: OK
test createtest: OK
test openiput: OK
test exitiput: OK
test iput: OK
test mem: OK
test pipe1: OK
test preempt: kill... wait... OK
test exitwait: OK
test rmdot: OK
test fourteen: OK
test bigfile: OK
test dirfile: OK
test iref: OK
test forktest: OK
test bigdir: OK
ALL TESTS PASSED
$
```

3.Xv6 lab: Lock/Buffer cache

3.1 题目要求

测试命令 `bcachetest` 创建多个进程，这些进程重复读取不同的文件，如果多个进程密集使用文件系统，则它们可能会争用 `bcache.lock`。所以输出 `bcache` 锁的获取循环迭代次数将很高。题目要求修改块缓存，以便在运行 `bcachetest` 时，`bcache` 中所有锁的获取循环迭代次数接近于零。题目建议使用一个哈希表在缓存中查找块号，该哈希表每个哈希存储桶均具有锁定状态。

3.2 实验代码

1. 选取题目建议的 `bucket` 大小 13，使用固定数量的存储桶，不动态调整哈希表的大小。

```
1. #define NBUCKET 13
2. struct {
3.     struct spinlock lock[NBUCKET]; //每个 bucket 都有一个锁
4.     struct buf buf[NBUF];
5.
6.     // Linked list of all buffers, through prev/next.
7.     // Sorted by how recently the buffer was used.
8.     // head.next is most recent, head.prev is least.
9.     struct buf hashbucket[NBUCKET]; //哈希表，数据结构为链表
10. } bcache;
11.
12.
13. uint idx(uint blockno) //散列函数
14. {
15.     return blockno % NBUCKET;
16. }
```

2. 在 `kernel/buf.h` 的 `struct` 中添加字段 `time_stamp`，用以标记 `buf` 的时间戳。通过此更改，`brelse` 不需要获取 `bcache` 锁，并且 `bget` 可以根据时间戳选择最近最少使用的块。

```
1. struct buf {
2.     int valid; // has data been read from disk?
3.     int disk; // does disk "own" buf?
4.     uint dev;
5.     uint blockno;
6.     struct sleeplock lock;
7.     uint refcnt;
8.     struct buf *prev; // LRU cache list //最少使用的
9.     struct buf *next; // 最近使用的
10.     uchar data[BSIZE];
11.     uint time_stamp; //时间戳
12. };
```

3. 修改 kernel/bio.c 中的 binit() 函数。

```
1. void
2. binit(void)
3. {
4.     struct buf *b;
5.     for (int i=0; i<NBUCKETS;i++){
6.         initlock(&bcache.lock[i], "bcache");// 将 bucket 的头节点初始化为自己
7.         bcache.hashbucket[i].prev = &bcache.hashbucket[i];
8.         bcache.hashbucket[i].next = &bcache.hashbucket[i];
9.     }
10.
11.     for(b = bcache.buf; b < bcache.buf+NBUF; b++){
12.         b->time_stamp = ticks; // 记录一下 b 的时间戳
13.         b->next = bcache.hashbucket[0].next;
14.         b->prev = &bcache.hashbucket[0];
15.         initsleeplock(&b->lock, "buffer");
16.         bcache.hashbucket[0].next->prev = b;
17.         bcache.hashbucket[0].next = b;
18.     }
19. }
```

4. 修改 bget 函数，在哈希表中搜索缓冲区，并在找不到缓冲区时为该缓冲区分配一个条目，这必须是原子的。

```
1. static struct buf*
2. bget(uint dev, uint blockno)
3. {
4.     struct buf *b;
5.     int i= idx(blockno);//获取当前缓冲区块号的哈希码
6.     acquire(&bcache[i].lock);
7.
8.     // Is the block already cached?
9.     for(b = bcache.hashbucket[i].next; b != &bcache.hashbucket[i]; b = b->next){
10.         if(b->dev==dev && b->blockno == blockno)//hit
11.         {
12.             b->time_stamp = ticks; // 记录一下 b 的时间戳
13.             b->refcnt++;
14.             release(&bcache.lock[i]);
15.             acquiresleep(&b->lock);
16.             return b;
17.         } // 找到最小的时间戳对应 buf
```

```

18.
19. }
20.
21. // Not cached.如果找不到，就需要从其它 bucket 中寻找
22. // Recycle the least recently used (LRU) unused buffer.
23. for (int newnum = idx(blockno+1); newnum != i; newnum = (newnum+1)%NBUCKET
    ){
24.     acquire(&bcache.lock[newnum]);
25.     // 遍历该 bucket 的链表
26.     for(b = bcache.hashbucket[newnum].prev; b != &bcache.hashbucket[newnum]; b = b->prev){
27.         if(b->refcnt == 0) { // 找到了，修改该缓冲区块相应属性，并移到对应的
            bucket 中
28.             b->time_stamp = ticks; // 记录一下 b 的时间戳
29.             b->dev = dev;
30.             b->blockno = blockno;
31.             b->valid = 0;
32.             b->refcnt = 1;
33.
34.             // 先移出原来 bucket
35.             b->next->prev = b->prev;
36.             b->prev->next = b->next;
37.             release(&bcache.lock[newnum]);
38.
39.             //放到需要该缓冲区块的 bucket
40.             b->next = bcache.hashbucket[i].next;
41.             b->prev = &bcache.hashbucket[i];
42.
43.             //将该区块连接到当前 bucket 中
44.             bcache.hashbucket[i].next->prev = b;
45.             bcache.hashbucket[i].next = b;
46.             release(&bcache.lock[i]);
47.             acquiresleep(&b->lock);
48.             return b;
49.         }
50.     }
51.     release(&bcache.lock[newnum]); // 当前 bucket 中没找到，仍要释放锁
52. }
53. panic("bget: no buffers");
54. }

```

5. 修改 brelse, bpin, bunpin 函数。

```

1. void
2. brelse(struct buf *b)
3. {
4.     if(!holdingsleep(&b->lock))
5.         panic("brelse");
6.     releasesleep(&b->lock);
7.
8.     int i=idx(b->blockno);
9.
10.    b->time_stamp = ticks;//用时间戳代替原先的上锁
11.    if(b->time_stamp == ticks){
12.        b->refcnt--;
13.        if (b->refcnt == 0) {
14.            // no one is waiting for it.
15.            b->next->prev = b->prev;
16.            b->prev->next = b->next;
17.            b->next = bcache.hashbucket[i].next;
18.            b->prev = &bcache.hashbucket[i];
19.            bcache.hashbucket[i].next->prev = b;
20.            bcache.hashbucket[i].next = b;
21.        }
22.    }
23. }

```

```

1. void
2. bpin(struct buf *b) {
3.     int i=idx(b->blockno);
4.     acquire(&bcache.lock[i]);
5.     b->refcnt++;
6.     release(&bcache.lock[i]);
7. }
8.
9. void
10. bunpin(struct buf *b) {
11.     int i=idx(b->blockno);
12.     acquire(&bcache.lock[i]);
13.     b->refcnt--;
14.     release(&bcache.lock[i]);
15. }

```

3.3 实验结果

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
start test0
test0 results:
--- lock kmem/bcache stats
lock: kmem: #fetch-and-add 0 #acquire() 32901
lock: kmem: #fetch-and-add 0 #acquire() 61
lock: kmem: #fetch-and-add 0 #acquire() 84
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8392
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8698
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10442
lock: bcache.bucket: #fetch-and-add 0 #acquire() 10756
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8526
lock: bcache.bucket: #fetch-and-add 0 #acquire() 9898
lock: bcache.bucket: #fetch-and-add 0 #acquire() 8478
--- top 5 contended locks:
lock: virtio_disk: #fetch-and-add 9070749 #acquire() 1026
lock: proc: #fetch-and-add 1668070 #acquire() 90544
lock: proc: #fetch-and-add 1469935 #acquire() 90523
lock: proc: #fetch-and-add 1332378 #acquire() 90522
lock: proc: #fetch-and-add 1282240 #acquire() 90523
tot= 0
test0: OK
start test1
test1 OK
$ █
```

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
test0: OK
start test1
test1 OK
$ usertests
usertests starting
test manywrites: OK
test execout: OK
test copyin: OK
test copyout: OK
test copyinstr1: OK
test copyinstr2: OK
test copyinstr3: OK
test rwsbrk: OK
test truncate1: OK
test truncate2: OK
test truncate3: OK
test reparent2: OK
test pgbug: OK
test sbrkbugs: usertrap(): unexpected scause 0x000000000000000c pid=3252
                sepc=0x00000000000056a0 stval=0x00000000000056a0
usertrap(): unexpected scause 0x000000000000000c pid=3253
                sepc=0x00000000000056a0 stval=0x00000000000056a0
OK
test badarg: OK
```

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
test validatetest: OK
test stacktest: usertrap(): unexpected scause 0x000000000000000d pid=6288
      sepc=0x000000000000022cc stval=0x000000000000fb90
OK
test opentest: OK
test writetest: OK
test writebig: OK
test createtest: OK
test openinput: OK
test exitinput: OK
test input: OK
test mem: OK
test pipe1: OK
test preempt: kill... wait... OK
test exitwait: OK
test rmdot: OK
test fourteen: OK
test bigfile: OK
test dirfile: OK
test iref: OK
test forktest: OK
test bigdir: OK
ALL TESTS PASSED
$
```

4.Xv6 lab: File System/Large files

4.1 题目要求

测试命令 `bigfile` 希望能够创建具有 65803 个块的文件，但是未经修改的 `xv6` 会将文件限制为 268 个块。这是因为：`xv6` 索引节点包含 12 个“直接”块编号和一个“单间接”块编号，这是指最多容纳 256 个以上块编号的块，总共 $12 + 256 = 268$ 块。

所以需要更改 `xv6` 文件系统代码，以在每个 `inode` 中支持“双重间接”块，其中包含 256 个单间接块地址，每个间接块最多可以包含 256 个数据块地址。结果将是一个文件最多可以包含 65803 个块或 $256 * 256 + 256 + 11$ 个块（11 个代替 12 个，因为我们将为双间接块牺牲一个直接块号）。

4.2 实验代码

1. 修改 `fs.h` 中的宏定义以及 `dinode` 结构，为其添加一个双重间接指针，同时将直接指针-1, `addrs[0-10]` 对应 11 个直接指针，`addrs[11]` 对应间接指针，`addrs[12]` 对应双重间接指针。

```

#define NDIRECT 11 //direct pointer 11
#define NINDIRECT1 (BSIZE / sizeof(uint)) //indirect pointer 256
#define NINDIRECT2 (NINDIRECT1 * NINDIRECT1) //double indirect pointer 256*256
#define NINDIRECT (NINDIRECT1 + NINDIRECT2)
#define MAXFILE (NDIRECT + NINDIRECT) //size of file

// On-disk inode structure
struct dinode {
    short type; // File type
    short major; // Major device number (T_DEVICE only)
    short minor; // Minor device number (T_DEVICE only)
    short nlink; // Number of links to inode in file system
    uint size; // Size of file (bytes)
    //addr[0-10] are direct pointer, addr[1] is indirect pointer,
    //addr[12] is double indirect pointer
    uint addr[NDIRECT+2]; // Data block addresses
};

```

2. 然后修改 fs.c 中的 bmap() 函数。

```

1. static uint
2. bmap(struct inode *ip, uint bn)
3. {
4.     uint addr, *a;
5.     struct buf *bp;
6.
7.     if(bn < NDIRECT){ //bn 仍在直接指针范围内，可以直接访问
8.         if((addr = ip->addrs[bn]) == 0)
9.             ip->addrs[bn] = addr = balloc(ip->dev);
10.        return addr;
11.    }
12.
13.    bn -= NDIRECT; //减去直接指针，判断在不在间接指针范围内
14.    if(bn < NINDIRECT1){
15.        // Load indirect block, allocating if necessary.
16.        if((addr = ip->addrs[NDIRECT]) == 0) //访问间接指针，没有就分配一个
17.            ip->addrs[NDIRECT] = addr = balloc(ip->dev);
18.        bp = bread(ip->dev, addr); //访问间接指针指向的指针
19.        a = (uint*)bp->data; //访问间接指针指向的指针所指的内容
20.        if((addr = a[bn]) == 0){ //为空就分配一个
21.            a[bn] = addr = balloc(ip->dev);
22.            log_write(bp);
23.        }
24.        brelse(bp);
25.        return addr;
26.    }
27.
28.    bn -= NINDIRECT1; //减去间接指针，判断在不在双重间接指针范围内
29.    if(bn < NINDIRECT2){

```



```

30.     uint bn_1= ( bn & 0xff00)>>8 ;//双重间接指针中的一级指针
31.     uint bn_2= bn & 0xff;//双重间接指针中的二级指针
32.
33.     // Load double indirect block, allocating if necessary.
34.     if((addr = ip->addrs[NDIRECT+1]) == 0)//访问双重间接指针，没有就分配一个
35.         ip->addrs[NDIRECT+1] = addr = balloc(ip->dev);
36.     bp = bread(ip->dev, addr);
37.     a = (uint*)bp->data;
38.     if((addr = a[bn_1]) == 0){//访问一级指针
39.         a[bn_1] = addr = balloc(ip->dev);
40.         log_write(bp);
41.     }
42.     brelse(bp);
43.
44.     bp = bread(ip->dev, addr);
45.     a = (uint*)bp->data;
46.     if((addr = a[bn_2]) == 0){//访问二级指针
47.         a[bn_2] = addr = balloc(ip->dev);
48.         log_write(bp);
49.     }
50.     brelse(bp);
51.
52.     return addr;
53. }
54.
55. panic("bmap: out of range");
56. }

```

3. 修改 itrunc 函数，主要功能是释放所有的指针。

```

1. void
2. itrunc(struct inode *ip)
3. {
4.     int i, j, k;
5.     struct buf *bp;
6.     struct buf *bp2;
7.     uint *a;
8.     uint *a2;
9.
10.    for(i = 0; i < NDIRECT; i++){//释放直接指针
11.        if(ip->addrs[i]){
12.            bfree(ip->dev, ip->addrs[i]);
13.            ip->addrs[i] = 0;
14.        }
15.    }

```

```

16.
17.  if(ip->addrs[NDIRECT]){//释放间接指针
18.      bp = bread(ip->dev, ip->addrs[NDIRECT]);
19.      a = (uint*)bp->data;
20.      for(j = 0; j < NINDIRECT1; j++){
21.          if(a[j])
22.              bfree(ip->dev, a[j]);
23.      }
24.      brelse(bp);
25.      bfree(ip->dev, ip->addrs[NDIRECT]);
26.      ip->addrs[NDIRECT] = 0;
27.  }
28.
29.  if(ip->addrs[NDIRECT+1]){//释放双重间接指针
30.      bp = bread(ip->dev, ip->addrs[NDIRECT+1]);
31.      a = (uint*)bp->data;
32.      for(j = 0; j < NINDIRECT1; j++){//释放一级指针
33.          if(a[j])
34.          {
35.              bp2 = bread(ip->dev, a[j]);
36.              a2 = (uint*)bp2->data;
37.              for(k=0 ;k < NINDIRECT1 ;k++){//释放二级指针
38.                  {
39.                      if(a2[k])
40.                          bfree(ip->dev,a2[k]);
41.                  }
42.                  brelse(bp2);
43.                  bfree(ip->dev, a[j]);
44.                  a[j]=0;
45.              }
46.          }
47.          brelse(bp);
48.          bfree(ip->dev, ip->addrs[NDIRECT+1]);
49.          ip->addrs[NDIRECT+1] = 0;
50.      }
51.
52.      ip->size = 0;
53.      iupdate(ip);
54. }

```

4.3 实验结果

实验结果如图，可以看出已正确输出预期结果。

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
xv6 kernel is booting

init: starting sh
$ bigfile
.....
wrote 65803 blocks
bigfile done; ok
$ usertests
usertests starting
test manywrites: OK
test execout: OK
test copyin: OK
test copyout: OK
test copyinstr1: OK
test copyinstr2: OK
test copyinstr3: OK
```

.....

.....

```
niandd33@ubuntu: ~/xv6-labs-2020
File Edit View Search Terminal Help
test validatetest: OK
test stacktest: usertrap(): unexpected scause 0x000000000000000d pid=6219
      sepc=0x000000000000022ce stval=0x000000000000fb90
OK
test opentest: OK
test writetest: OK
test writebig: OK
test createtest: OK
test openinput: OK
test exitinput: OK
test input: OK
test mem: OK
test pipe1: OK
test preempt: kill... wait... OK
test exitwait: OK
test rmdot: OK
test fourteen: OK
test bigfile: OK
test dirfile: OK
test iref: OK
test forktest: OK
test bigdir: OK
ALL TESTS PASSED
$
```